



LinkedIn

Junsu Lee

junsulee.tech@gmail.com | [linkedin.com/in/junsuleetech](https://www.linkedin.com/in/junsuleetech) | github.com/YieumYoon | yieumyoon.github.io | Cleveland, OH | Willing to relocate nationwide



GitHub

PROFESSIONAL SUMMARY

Computer Science senior graduating December 2026. Deployed a Boston Dynamics Spot inspection program in a live factory, fine-tuned and physically evaluated a bimanual vision-language-action model, and connected Codex/OpenCodex to a monitored local Qwen 27B serving stack.

EDUCATION

Bachelor of Arts in Computer Science | GPA: 3.47 Case Western Reserve University
Cleveland, OH

Expected December 2026

Coursework: Intro Artificial Intelligence, Fundamentals of Robotics, Computational Perception, Software Engineering. President, CWRU ACM Student Chapter (August 2026 - Present).

RELEVANT EXPERIENCE

Bendix Commercial Vehicle Systems

January 2026 - Present

Advanced Manufacturing Engineer Co-op

Avon, OH

- Took Boston Dynamics Spot inspection from manual setup to production operation in a live factory. Configured approximately 500 inspection points, moving plant-wide collection from roughly every six months to daily or every-other-day autonomous operation.
- Shipped eight field-used tools for Orbit run-history review, live sensor diagnostics, map validation, worker signaling, and alerts. Built map and inspection-point editing workflows that cut a full update from at least one person-week to within one workday while retaining human approval.
- Built image- and log-rich alerts that cut incident discovery and response from as long as 12-24 hours to approximately 15-30 minutes; integrated Spot, GraphNav, Orbit APIs, Protocol Buffers, camera, depth, and robot-state data.

xLab at Case Western Reserve University | Progressive Insurance partner

June 2025 - August 2025

Weather AI Research Intern

Cleveland, OH

- Owned model execution and backend delivery on a three-student R&D team. Built a Python Earth2Studio/FourCastNet pipeline on six A6000 GPUs, generating seven-day forecasts in about three minutes and publishing NetCDF4 and preview outputs to AWS S3.

xLab at Case Western Reserve University | Eaton partner

January 2025 - May 2025

Frontend Developer, AI Call Center Training System

Cleveland, OH

- Built Next.js/TypeScript frontend components and integrated FastAPI API calls for an AI customer-simulation proof of concept designed to help call-center agents practice technical conversations with varied emotions and multilingual scenarios.

Republic of Korea Army

January 2023 - July 2024

Squad Leader / IT Technician (Signaler)

South Korea

- Improved CCTV field installation through bench-testing and preassembly, reducing on-site installation time by approximately 70%; trained a five-person team supporting more than 400 networked devices and 34 CCTV units while troubleshooting LAN and communications equipment.

PROJECTS

Language-Guided Arm Selection on a Bimanual SO-101 with NVIDIA GR00T N1.5

Owned the core pipeline for a custom bimanual SO-101 embodiment, from three-camera teleoperation and dataset repair to training and physical evaluation. Re-recorded approximately 360 validated episodes totaling about 238,000 frames after diagnosing timestamp corruption. Fine-tuned GR00T N1.5 on H200/L40S GPUs, cut projected training time from about 15 hours to 2.5 hours, and reached approximately 25% full-task success across roughly 20 physical trials versus 0% untuned.

Local Owen Inference and Observability System

Operate an always-on Ubuntu/RTX 5090 server for local Qwen inference via Codex/OpenCodex, llama-swap, and vLLM; AI handled most configuration after OS installation. Switch vLLM, ComfyUI, and WhisperX with one GPU-resident service at a time; ComfyUI/WhisperX run in Docker. Services restart after reboot; Tailscale enables remote access. Monitor GPU temperature/utilization/power, vLLM throughput/requests/KV cache, and CPU/RAM/disk in Grafana.

Spot Fiducial AprilTag Plate Generator

Built a public JavaScript/SVG tool for Spot localization tags; automated SVG-to-STEP conversion through Fusion 360, cutting per-tag CAD preparation from as much as two hours to roughly 10 seconds across more than 120 plates.

Data Pipelines and Visualization

Built Python/SQL ELT workflows and six tested dbt analytical models; orchestrated an Airflow/PostgreSQL environmental-data pipeline with Tableau dashboards.

SKILLS

Languages: Python, TypeScript/Javascript, Java, SQL, Bash; C/C++ (coursework)

AI & ML: PyTorch, TensorFlow, OpenCV, Hugging Face, Weights & Biases, NumPy

Robotics & Physical AI: Boston Dynamics Spot, GraphNav, Orbit APIs, SO-101, GR00T N1.5, LeRobot, Foxglove

LLM Systems & GenAI: vLLM, llama-swap, Qwen, Grafana; Gemini API, RAG, function calling (Google/Kaggle labs)

Software & Data: React, Next.js, FastAPI (API integration), Pandas, dbt, Airflow, PostgreSQL

Infrastructure & Tools: Git, Docker/Podman, Linux/HPC, AWS S3, Makefile, tmux

CAD & Fabrication: Fusion 360, SolidWorks (CSWA), 3D printing

IT & Networking: LAN/switch diagnostics, CCTV, IP telephony

CERTIFICATIONS & TRAINING

NVIDIA-Certified Associate: Generative AI LLMs NVIDIA

Valid through October 2027

5-Day Generative AI Intensive (hands-on labs) Google & Kaggle

November 2024

Deep Learning Specialization DeepLearning.AI

March 2024